

# SAMPLE EFFICIENT ACTOR-CRITIC WITH EXPERIENCE REPLAY

**Ziyu Wang**  
DeepMind  
ziyu@google.com

**Victor Bapst**  
DeepMind  
vbapst@google.com

**Nicolas Heess**  
DeepMind  
heess@google.com

**Volodymyr Mnih**  
DeepMind  
vmnih@google.com

**Remi Munos**  
DeepMind  
Munos@google.com

**Koray Kavukcuoglu**  
DeepMind  
korayk@google.com

**Nando de Freitas**  
DeepMind, CIFAR, Oxford University  
nandodef Freitas@google.com

## 【强化学习算法 6】ACER



张楚琦

清华大学 交叉信息院博士在读

13 人赞同了该文章

从文章的标题可以看出ACER指的actor-critic with experience replay。

原文传送门：

Mnih, Volodymyr, et al. "Asynchronous methods for deep reinforcement learning." *International conference on machine learning*. 2016. (前序工作)

Munos, Rémi, et al. "Safe and efficient off-policy reinforcement learning." *Advances in Neural Information Processing Systems*. 2016. (前序工作)

Degrís, Thomas, Martha White, and Richard S. Sutton. "Off-policy actor-critic." *arXiv preprint arXiv:1205.4839* (2012). (前序工作)

Wang, Ziyu, et al. "Sample efficient actor-critic with experience replay." *arXiv preprint arXiv:1611.01224* (2016).

**特色：**强化学习里面和环境的交互成本是比较高的，我们希望一个算法能够尽可能少地与环境交互，这个特性称之为sample efficient。提高sample efficiency的一个好办法是使用experience replay，这样就需要使用off-policy的方法来做了。这里就展示了一种off-policy actor-critic方法。

**分类：**Model-free、**Policy-based**、**Off-policy**、Continuous State Space、Continuous Action Space (also discrete)、Support High-dim Input

**理论依据：**

对于off-policy的策略来说，拿到的轨迹和on-policy的轨迹不一样，需要使用importance ratio来对其权重进行调整之后才好使用。因此policy gradient可以被写作

$$\hat{g} = \left( \prod_{t=0}^h \rho_t \right) \sum_{t=0}^h \sum_{a_t} \gamma^t r_{t+1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

importance ratio的连乘是对于整个轨迹来做的，因子  $\prod_{t=0}^h \rho_t$  就很容易过大或者为零。Off-policy Actor-critic这篇工作说off-policy policy gradient大致上可以单步拆成这样

$$g = \mathbb{E}_{\rho} [\rho_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q^{\pi}(s_t, a_t)] \quad (1)$$

**过程：**

这篇工作下面就是围绕(1)式展开来做的。

1. 由于  $\rho_t = \frac{\pi(a_t|x_t)}{\mu(a_t|x_t)}$ ，这个数值很容易过大，使算法不稳定，因此使用了一个叫做importance weight truncation with bias correction的技术，做如下变换  $\mathbb{E}_\pi[\rho_t \dots] = \mathbb{E}_\pi[\tilde{\rho}_t \dots] + \mathbb{E}_{\pi \sim \pi} \left[ \left[ \frac{\rho_t(a) - c}{\rho_t(a)} \right]_+, \dots \right]$ ，其中  $\tilde{\rho}_t = \min(c, \rho_t)$ 。这样的变换既不会产生额外的bias，而且产生的两项单独都是bounded的，前一个小于c，后一个小于1。

2.  $Q^*(a_t, a_t)$  的估计使用了Retrace这种技术

$$Q^{\text{ret}}(x_t, a_t) = r_t + \gamma \bar{\rho}_{t+1} [Q^{\text{ret}}(x_{t+1}, a_{t+1}) - Q(x_{t+1}, a_{t+1})] + \gamma V(x_{t+1}),$$

3. 上述的Q和V的估计使用了dueling network的结构，对于连续的行动空间，文章提出了stochastic dueling network。主要是因为行动空间连续的时候之前dueling network的  $\sum_i A_{\theta_i}(x, a)$  没法算了，因此这里用采样的方法来计算。

$$\tilde{Q}_{\theta_v}(x_t, a_t) \sim V_{\theta_v}(x_t) + A_{\theta_v}(x_t, a_t) - \frac{1}{n} \sum_{i=1}^n A_{\theta_v}(x_t, u_i), \text{ and } u_i \sim \pi_\theta(\cdot|x_t)$$

网络输出的是  $Q_{\theta}$  和  $A_{\theta}$ ，然后通过这个方法拼起来得到Q。

4. 通过综合了以上技术，就得到了ACER的off-policy policy gradient

$$\begin{aligned} \hat{g}_t^{\text{acer}} &= \bar{\rho}_t \nabla_{\phi_\theta(x_t)} \log f(a_t|\phi_\theta(x_t)) [Q^{\text{ret}}(x_t, a_t) - V_{\theta_v}(x_t)] \\ &\quad + \mathbb{E}_{a \sim \pi} \left( \left[ \frac{\rho_t(a) - c}{\rho_t(a)} \right]_+ \nabla_{\phi_\theta(x_t)} \log f(a_t|\phi_\theta(x_t)) [Q_{\theta_v}(x_t, a) - V_{\theta_v}(x_t)] \right). \end{aligned}$$

下面需要确定每步更新的步长。希望步长产生的变化在策略空间不要太大，因此我们希望限制KL散度的变化，为了避免像TRPO那样对于KL关于  $\theta$  求Hessian，这里1) 使用了一个average policy network，其参数更新方式为  $\theta_a \leftarrow \alpha \theta_a + (1 - \alpha) \theta$ ，这样KL散度的变化主要是被一个一阶项主导，避免了求Hessian的问题；2) 都对于策略分布的参数  $\phi_\theta$  求导，所有的优化都在对  $\phi_\theta$  的梯度上算清楚了，再一次BP传递到  $\theta$  上，避免了多次BP经过神经网络的复杂运算。

通过写出trust region optimization problem

$$\begin{aligned} \underset{z}{\text{minimize}} \quad & \frac{1}{2} \|\hat{g}_t^{\text{acer}} - z\|_2^2 \\ \text{subject to} \quad & \nabla_{\phi_\theta(x_t)} D_{KL} [f(\cdot|\phi_{\theta_a}(x_t)) \| f(\cdot|\phi_\theta(x_t))]^T z \leq \delta \end{aligned}$$

直接解析求得最优解

$$z^* = \hat{g}_t^{\text{acer}} - \max \left\{ 0, \frac{k^T \hat{g}_t^{\text{acer}} - \delta}{\|k\|_2^2} \right\} k$$

可以得到参数更新公式

$$\theta \leftarrow \theta + \frac{\partial \phi_\theta(\pi)}{\partial \theta} z^*$$

算法:

---

**Algorithm 1** ACER for discrete actions (master algorithm)

---

// Assume global shared parameter vectors  $\theta$  and  $\theta_v$ .  
 // Assume ratio of replay  $r$ .

**repeat**

  Call ACER on-policy, Algorithm 2

$n \leftarrow \text{Possion}(r)$

**for**  $i \in \{1, \dots, n\}$  **do**

    Call ACER off-policy, Algorithm 2

**end for**

**until** Max iteration or time reached.

---



---

**Algorithm 2** ACER for discrete actions

---

Reset gradients  $d\theta \leftarrow 0$  and  $d\theta_v \leftarrow 0$ .

Initialize parameters  $\theta' \leftarrow \theta$  and  $\theta'_v \leftarrow \theta_v$ .

**if not** On-Policy **then**

  Sample the trajectory  $\{x_0, a_0, r_0, \mu(\cdot|x_0), \dots, x_k, a_k, r_k, \mu(\cdot|x_k)\}$  from the replay memory.

**else**

  Get state  $x_0$

**end if**

**for**  $i \in \{0, \dots, k\}$  **do**

  Compute  $f(\cdot|\phi_{\theta'}(x_i))$ ,  $Q_{\theta'_v}(x_i, \cdot)$  and  $f(\cdot|\phi_{\theta_a}(x_i))$ .

**if** On-Policy **then**

    Perform  $a_i$  according to  $f(\cdot|\phi_{\theta'}(x_i))$

    Receive reward  $r_i$  and new state  $x_{i+1}$

$\mu(\cdot|x_i) \leftarrow f(\cdot|\phi_{\theta'}(x_i))$

**end if**

$\bar{\rho}_i \leftarrow \min \left\{ 1, \frac{f(a_i|\phi_{\theta'}(x_i))}{\mu(a_i|x_i)} \right\}$ .

**end for**

$Q^{ret} \leftarrow \begin{cases} 0 & \text{for terminal } x_k \\ \sum_a Q_{\theta'_v}(x_k, a) f(a|\phi_{\theta'}(x_k)) & \text{otherwise} \end{cases}$

**for**  $i \in \{k-1, \dots, 0\}$  **do**

$Q^{ret} \leftarrow r_i + \gamma Q^{ret}$

$V_i \leftarrow \sum_a Q_{\theta'_v}(x_i, a) f(a|\phi_{\theta'}(x_i))$

  Computing quantities needed for trust region updating:

$$\begin{aligned} g &\leftarrow \min \{c, \rho_i(a_i)\} \nabla_{\phi_{\theta'}(x_i)} \log f(a_i|\phi_{\theta'}(x_i)) (Q^{ret} - V_i) \\ &\quad + \sum_a \left[ 1 - \frac{c}{\rho_i(a)} \right]_+ f(a|\phi_{\theta'}(x_i)) \nabla_{\phi_{\theta'}(x_i)} \log f(a|\phi_{\theta'}(x_i)) (Q_{\theta'_v}(x_i, a) - V_i) \\ k &\leftarrow \nabla_{\phi_{\theta'}(x_i)} D_{KL} [f(\cdot|\phi_{\theta_a}(x_i)) \| f(\cdot|\phi_{\theta'}(x_i))] \end{aligned}$$

Accumulate gradients wrt  $\theta'$ :  $d\theta' \leftarrow d\theta' + \frac{\partial \phi_{\theta'}(x_i)}{\partial \theta'} \left( g - \max \left\{ 0, \frac{k^T g - \delta}{\|k\|_2^2} \right\} k \right)$

Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \nabla_{\theta'_v} (Q^{ret} - Q_{\theta'_v}(x_i, a))^2$

Update Retrace target:  $Q^{ret} \leftarrow \bar{\rho}_i (Q^{ret} - Q_{\theta'_v}(x_i, a)) + V_i$

**end for**

Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .

Updating the average policy network:  $\theta_a \leftarrow \alpha \theta_a + (1 - \alpha) \theta$

---

知乎 @张楚衍



**Algorithm 3** ACER for Continuous Actions

---

Reset gradients  $d\theta \leftarrow 0$  and  $d\theta_v \leftarrow 0$ .  
Initialize parameters  $\theta' \leftarrow \theta$  and  $\theta'_v \leftarrow \theta_v$ .  
Sample the trajectory  $\{x_0, a_0, r_0, \mu(\cdot|x_0), \dots, x_k, a_k, r_k, \mu(\cdot|x_k)\}$  from the replay memory.  
**for**  $i \in \{0, \dots, k\}$  **do**  
    Compute  $f(\cdot|\phi_{\theta'}(x_i))$ ,  $V_{\theta'_v}(x_i)$ ,  $\tilde{Q}_{\theta'_v}(x_i, a_i)$ , and  $f(\cdot|\phi_{\theta_a}(x_i))$ .  
    Sample  $a'_i \sim f(\cdot|\phi_{\theta'}(x_i))$   
     $\rho_i \leftarrow \frac{f(a_i|\phi_{\theta'}(x_i))}{\mu(a_i|x_i)}$  and  $\rho'_i \leftarrow \frac{f(a'_i|\phi_{\theta'}(x_i))}{\mu(a'_i|x_i)}$   
     $c_i \leftarrow \min \left\{ 1, (\rho_i)^{\frac{1}{\alpha}} \right\}$ .  
**end for**  
 $Q^{ret} \leftarrow \begin{cases} 0 & \text{for terminal } x_k \\ V_{\theta'_v}(x_k) & \text{otherwise} \end{cases}$   
 $Q^{opc} \leftarrow Q^{ret}$   
**for**  $i \in \{k-1, \dots, 0\}$  **do**  
     $Q^{ret} \leftarrow r_i + \gamma Q^{ret}$   
     $Q^{opc} \leftarrow r_i + \gamma Q^{opc}$   
    Computing quantities needed for trust region updating:  

$$g \leftarrow \min \{c, \rho_i\} \nabla_{\phi_{\theta'}(x_i)} \log f(a_i|\phi_{\theta'}(x_i)) (Q^{opc}(x_i, a_i) - V_{\theta'_v}(x_i))$$

$$+ \left[ 1 - \frac{c}{\rho'_i} \right]_+ (\tilde{Q}_{\theta'_v}(x_i, a'_i) - V_{\theta'_v}(x_i)) \nabla_{\phi_{\theta'}(x_i)} \log f(a'_i|\phi_{\theta'}(x_i))$$

$$k \leftarrow \nabla_{\phi_{\theta'}(x_i)} D_{KL} [f(\cdot|\phi_{\theta_a}(x_i)) \| f(\cdot|\phi_{\theta'}(x_i))]$$
  
    Accumulate gradients wrt  $\theta$ :  $d\theta \leftarrow d\theta + \frac{\partial \phi_{\theta'}(x_i)}{\partial \theta'} \left( g - \max \left\{ 0, \frac{k^T g - \delta}{\|k\|_2^2} \right\} k \right)$   
    Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + (Q^{ret} - \tilde{Q}_{\theta'_v}(x_i, a_i)) \nabla_{\theta'_v} \tilde{Q}_{\theta'_v}(x_i, a_i)$   

$$d\theta_v \leftarrow d\theta_v + \min \{1, \rho_i\} \left( Q^{ret}(x_i, a_i) - \tilde{Q}_{\theta'_v}(x_i, a_i) \right) \nabla_{\theta'_v} V_{\theta'_v}(x_i)$$
  
    Update Retrace target:  $Q^{ret} \leftarrow c_i (Q^{ret} - \tilde{Q}_{\theta'_v}(x_i, a_i)) + V_{\theta'_v}(x_i)$   
    Update Retrace target:  $Q^{opc} \leftarrow (Q^{opc} - \tilde{Q}_{\theta'_v}(x_i, a_i)) + V_{\theta'_v}(x_i)$   
**end for**  
Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .  
Updating the average policy network:  $\theta_a \leftarrow \alpha \theta_a + (1 - \alpha) \theta$

---

知乎 @张楚珩

编辑于 2018-10-01

强化学习 (Reinforcement Learning)

算法

算法 (书籍)

赞同 13




9 条评论

分享

喜欢

收藏





强化学习前沿

读呀读paper

进入专栏